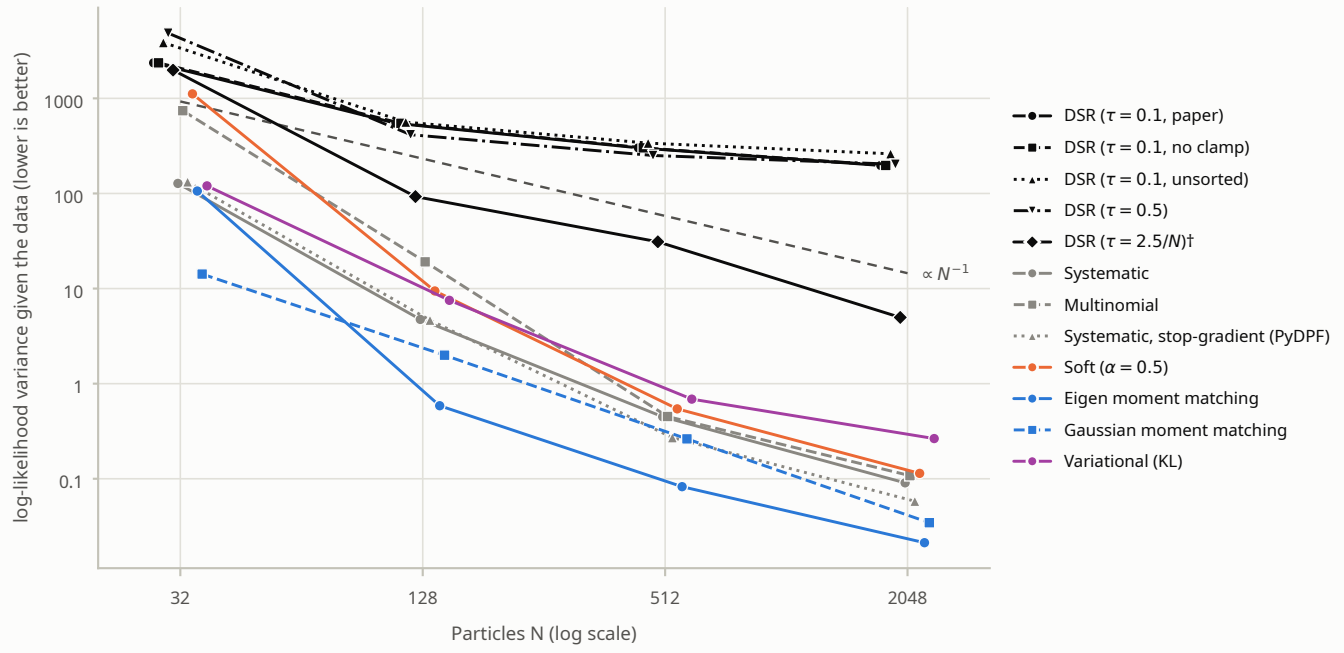
log-likelihood variance given the data (lower is better)
1000
100
10
1
0.1
32
128
512
2048
Particles N (log scale)
$\propto N^{-1}$
DSR ($\tau = 0.1$, paper)
DSR ($\tau = 0.1$, no clamp)
DSR ($\tau = 0.1$, unsorted)
DSR ($\tau = 0.5$)
DSR ($\tau = 2.5/N$)†
Systematic
Multinomial
Systematic, stop-gradient (PyDPF)
Soft ($\alpha = 0.5$)
Eigen moment matching
Gaussian moment matching
Variational (KL)