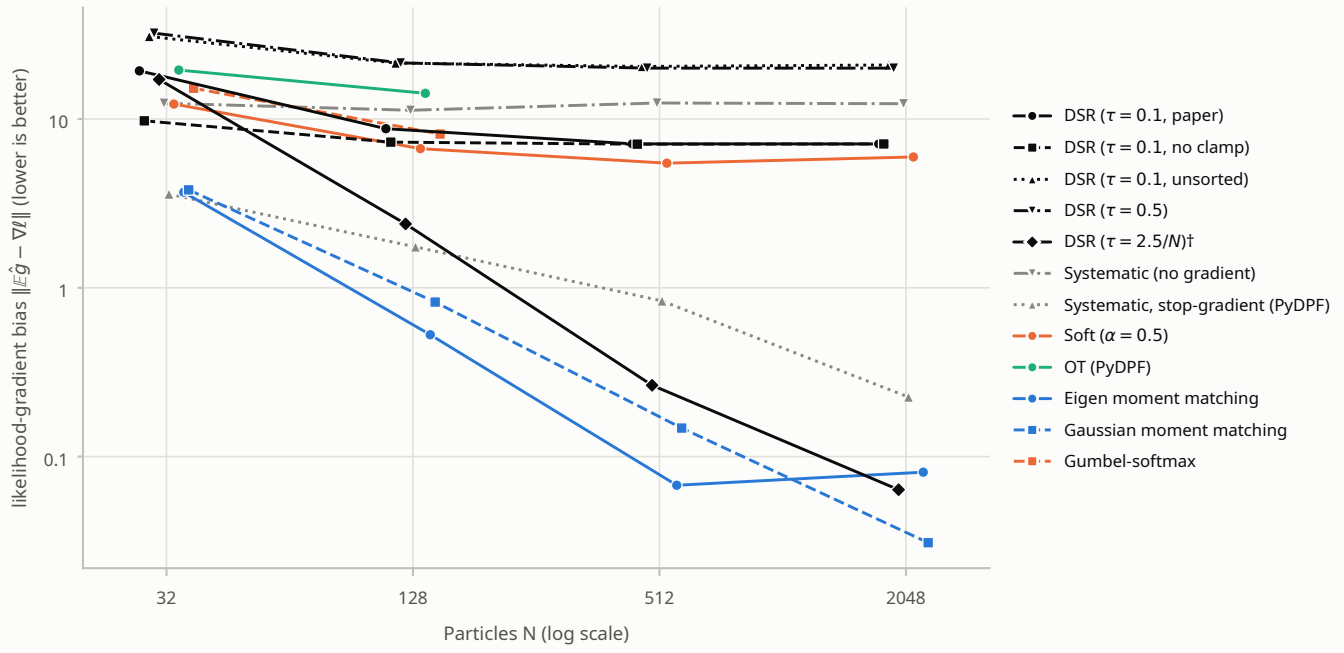
likelihood-gradient bias $\|\mathbb{E}\hat{g} - \nabla\ell\|$ (lower is better)
10
1
0.1
32
128
512
2048
Particles N (log scale)
DSR ($\tau = 0.1$, paper)
DSR ($\tau = 0.1$, no clamp)
DSR ($\tau = 0.1$, unsorted)
DSR ($\tau = 0.5$)
DSR ($\tau = 2.5/N$)†
Systematic (no gradient)
Systematic, stop-gradient (PyDPF)
Soft ($\alpha = 0.5$)
OT (PyDPF)
Eigen moment matching
Gaussian moment matching
Gumbel-softmax