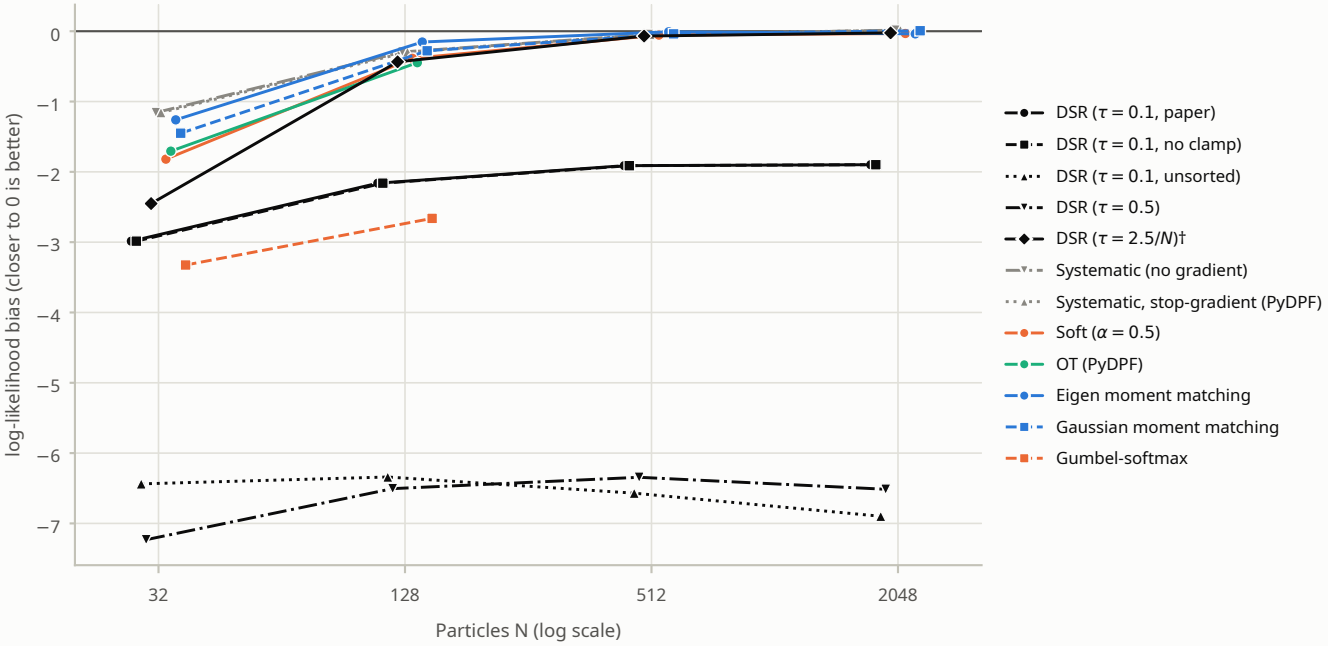
log-likelihood bias (closer to 0 is better)
0
−1
−2
−3
−4
−5
−6
−7
32
128
512
2048
Particles N (log scale)
DSR (τ = 0.1, paper)
DSR (τ = 0.1, no clamp)
DSR (τ = 0.1, unsorted)
DSR (τ = 0.5)
DSR (τ = 2.5/N)†
Systematic (no gradient)
Systematic, stop-gradient (PyDPF)
Soft (α = 0.5)
OT (PyDPF)
Eigen moment matching
Gaussian moment matching
Gumbel-softmax