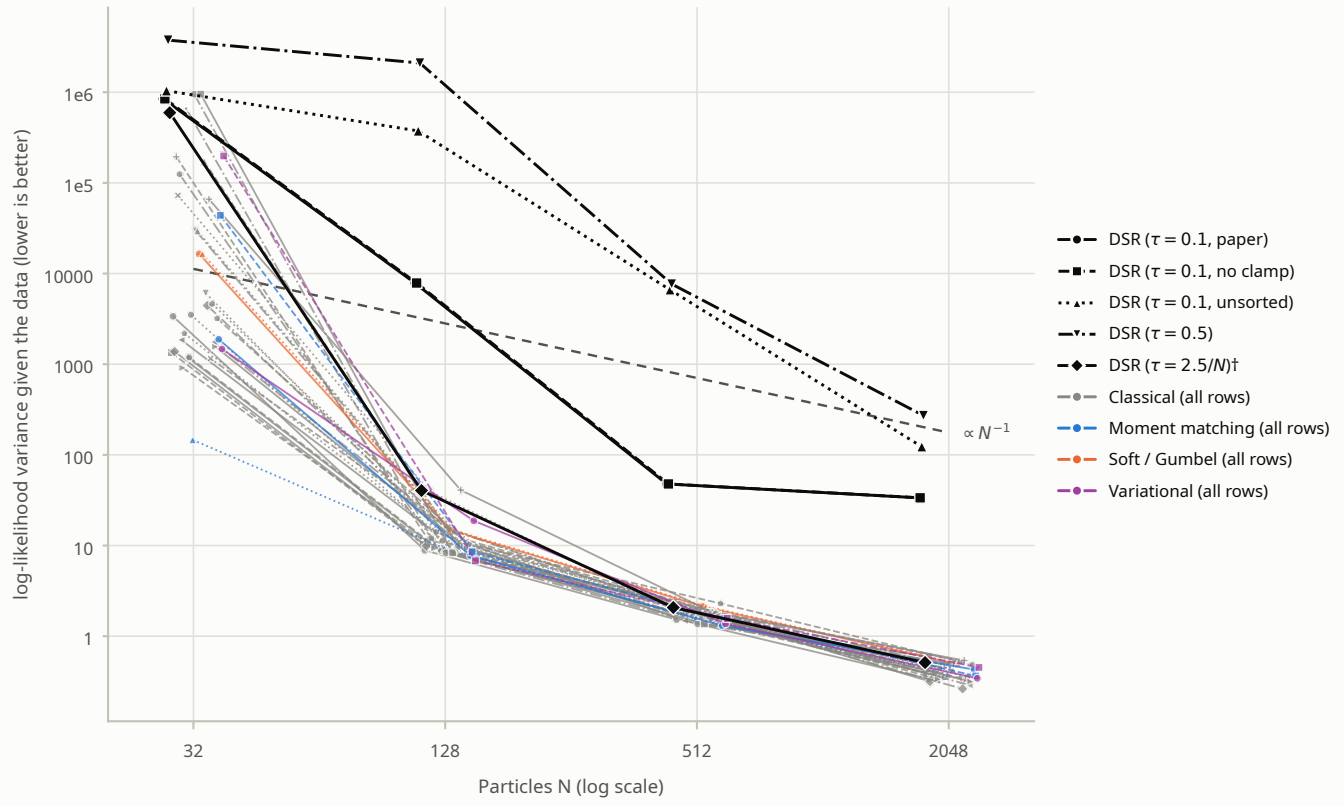
log-likelihood variance given the data (lower is better)
1e6
1e5
10000
1000
100
10
1
32
128
512
2048
Particles N (log scale)
∝ N^{-1}
DSR ($\tau = 0.1$, paper)
DSR ($\tau = 0.1$, no clamp)
DSR ($\tau = 0.1$, unsorted)
DSR ($\tau = 0.5$)
DSR ($\tau = 2.5/N$)†
Classical (all rows)
Moment matching (all rows)
Soft / Gumbel (all rows)
Variational (all rows)